

# Data challenges

---

- Non-standardized data
  - Inconveniently structured data
    - Tidy-ing data
    - Data with multiple factors
  - Duplicate data
  - Incorrect values
  - Missing values
- 



# Standardizing data

Raw Year	Standardized
2019	2019
'19	2019

Raw Medication	Standardized
azithromycin	azithromycin
Zithromax	azithromycin

Raw Name	Standardized
McDougal	McDougal
mcdougal	McDougal

Raw Unit	Standardized
micron	µm
µm	µm

# On dates

- ISO 8601 is an international standard for date-time information.
  - 20191001T182618+0000
- A challenge with this is that it requires dates be parsed again to do calculations.
  - It may be easier to store this as separate fields: year, month, day, hour, minute, seconds.
- Can use `pd.to_datetime` to convert to `DateTime` objects to allow subtraction, comparison.
- An alternative: Unix time
  - Number of seconds since 1 January 1970 UTC.
  - Returned by `time.time()` or e.g.

```
((pd.to_datetime('June 10, 2020') -  
  pd.to_datetime('January 1, 1970')) /  
  pd.Timedelta('1 sec'))
```



# Tidy Data

A data structuring approach.

Every variable has its own column.

Every observation has its own row.

Every value has its own cell.

# TODO: Melting

- Example and code



“8-mg Zofran”

Dosage (mg): 8  
Medicine: Zofran



“effusion of the right  
knee”

Condition: “knee effusion”  
Side: “right”



“white male”

Gender: male  
Ethnicity: Caucasian

Data with multiple factors

# Duplicate data

- Sometimes a data point (row) may be listed more than once, especially if manual entry was involved.
- But be careful: depending on how your data is structured, it may also be the case that data should appear more than once.
  - Imagine, e.g. a patient sent home from the hospital only to return later that day with the same conditions.

# Duplicate data example

```
import pandas as pd

patients = [
    (1002, 'Smith', 'John', 42, '20191001', 'diabetes'),
    (4261, 'Smith', 'Jane', 46, '20190510', 'pulmonary embolism'),
    (1002, 'Smith', 'John', 42, '20191001', 'diabetes'),
    (4171, 'Smith', 'Janet', 16, '20190909', 'acne')
]

data = pd.DataFrame(
    patients,
    columns=['pid', 'last', 'first', 'age', 'date', 'condition'])
```



# Duplicate data example

- See duplicate rows:

```
>>> data[data.duplicated()]
```

	pid	last	first	age	date	condition
2	1002	Smith	John	42	20191001	diabetes

# Duplicate data example

Pid should uniquely identify a patient.

Date and pid *almost* uniquely identifies an encounter.

- See duplicate rows:

```
>>> data[data.duplicated()]
```

	pid	last	first	age	date	condition
2	1002	Smith	John	42	20191001	diabetes



# Duplicate data example

- See duplicate rows:

```
>>> data[data.duplicated()]
      pid  last first  age      date condition
2  1002  Smith  John   42  20191001  diabetes
```

- Drop duplicate rows

```
>>> deduplicated_data = data.drop_duplicates()
>>> deduplicated_data
      pid  last first  age      date      condition
0  1002  Smith  John   42  20191001      diabetes
1  4261  Smith  Jane   46  20190510  pulmonary embolism
3  4171  Smith  Janet  16  20190909              acne
```

# Duplicate data example

- Both `data.duplicated` and `data.drop_duplicates` take an optional *subset* keyword argument specifying which columns to pay attention to.

```
>>> data.duplicated(subset=['last'])  
0      False  
1       True  
2       True  
3       True  
dtype: bool
```

# Incorrect values

- Define and check ranges
  - If a person is 57 years old, that is plausible. If a person is 577 years old, then maybe there is something wrong.
- Check categorical values
  - e.g. is the "State" field correct? We know the list of all possible states.
- Look for inconsistencies
  - e.g. City: "New Haven", Zip: "90210"
- Look at outliers
  - If only one person has a disease, it could be very rare... or it could be a typo.
- Validate when possible.